

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("/Users/bob/Downloads/11_winequality-red.csv")
df.fillna(0,inplace=True)
df
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66

1599 rows × 12 columns

```
In [3]: df.head()
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	

In [4]: `df.info()`

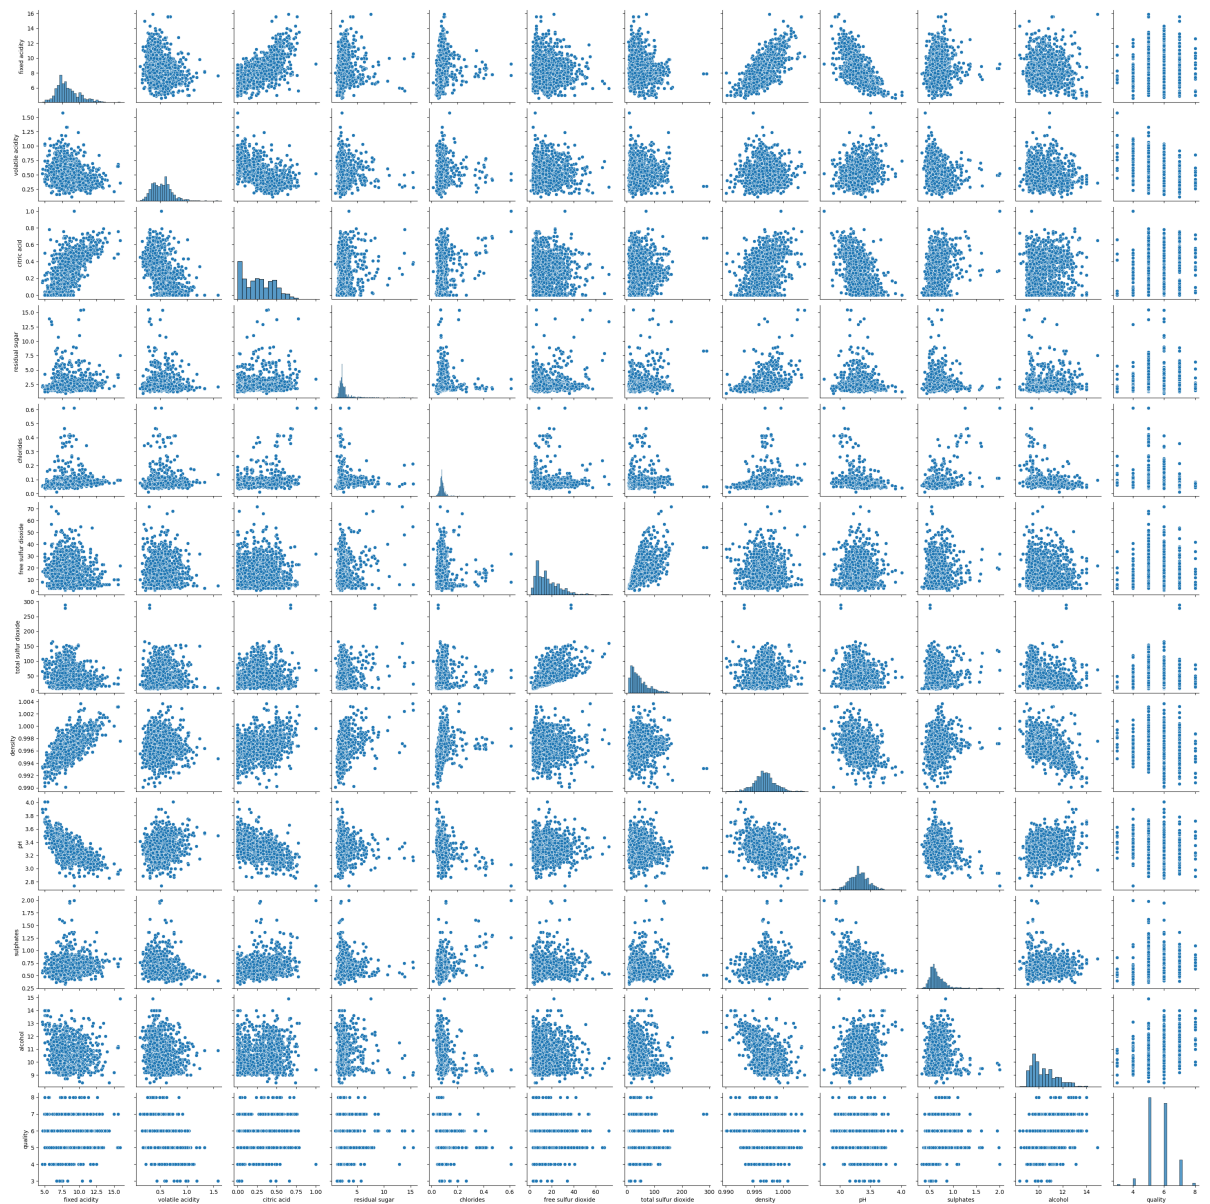
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [5]: `df.columns`

```
Out[5]: Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residu
al sugar',
              'chlorides', 'free sulfur dioxide', 'total sulfur dioxide',
              'density',
              'pH', 'sulphates', 'alcohol', 'quality'],
              dtype='object')
```


```
In [6]: sns.pairplot(df)
```

```
Out[6]: <seaborn.axisgrid.PairGrid at 0x7f8ef1bed900>
```



```
In [7]: sns.displot(df['quality'])
```

```
Out[7]: <seaborn.axisgrid.FacetGrid at 0x7f8ec91de980>
```

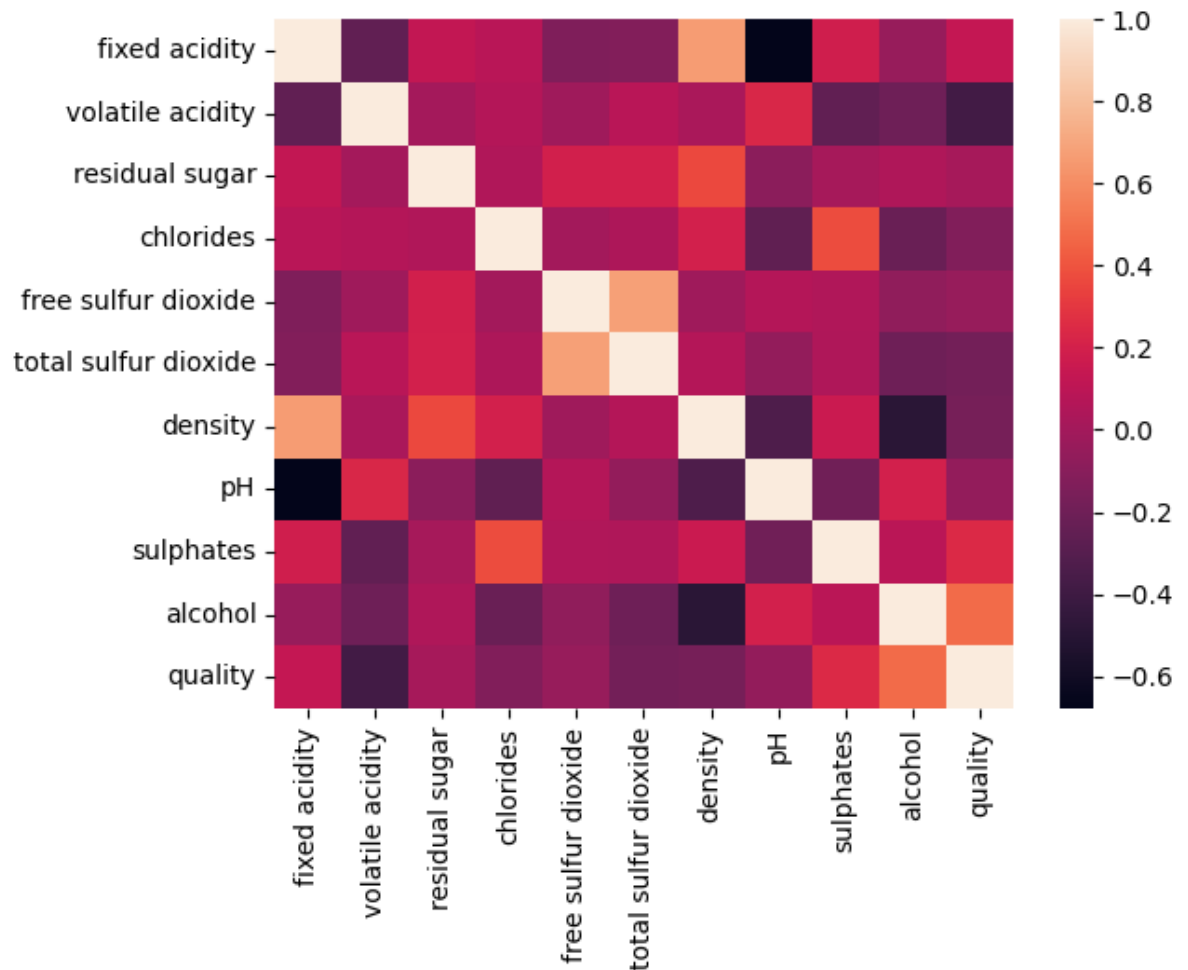


```
In [8]: df1=df.drop(['citric acid'],axis=1)
df1
df1=df1.drop(df1.index[1537:])
df1.isna().sum()
```

```
Out[8]: fixed acidity      0
volatile acidity    0
residual sugar      0
chlorides           0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
dtype: int64
```

```
In [9]: sns.heatmap(df1.corr())
```

```
Out[9]: <Axes: >
```



```
In [10]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [11]: df1.isna().sum()
```

```
Out[11]: fixed acidity      0
volatile acidity    0
residual sugar      0
chlorides           0
free sulfur dioxide 0
total sulfur dioxide 0
density             0
pH                  0
sulphates           0
alcohol             0
quality             0
dtype: int64
```

```
In [12]: y=df1['quality']
x=df1.drop(['volatile acidity'],axis=1)
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
print(x_train)
```

	fixed acidity	residual sugar	chlorides	free sulfur dioxid
e \				
792	7.1	2.5	0.081	17.
0				
974	8.8	5.9	0.073	7.
0				
901	7.4	2.4	0.080	16.
0				
1436	10.0	1.6	0.169	27.
0				
435	12.3	2.3	0.091	6.
0				
...	
...				
515	8.5	6.1	0.122	34.
0				
350	10.7	2.7	0.107	17.
0				
460	9.2	2.3	0.083	14.
0				
68	9.3	2.0	0.074	27.
0				
1497	6.9	2.3	0.054	7.
0				

	total sulfur dioxide	density	pH	sulphates	alcohol	qua
lity						
792	87.0	0.99745	3.48	0.60	9.7	
6						
974	13.0	0.99658	3.30	0.62	12.1	
7						
901	33.0	0.99736	3.58	0.69	10.8	
7						
1436	90.0	0.99914	3.15	0.65	8.5	
5						
435	18.0	1.00040	3.16	0.49	9.5	
5						
...	
...						
515	151.0	1.00100	3.31	1.14	9.3	
5						
350	34.0	1.00040	3.28	0.98	9.9	
6						
460	23.0	0.99760	3.35	0.61	11.3	
6						
68	65.0	0.99690	3.28	0.79	10.7	
5						
1497	16.0	0.99508	3.45	0.63	11.5	
6						

[1075 rows x 10 columns]

```
In [13]: model=LinearRegression()  
model.fit(x_train,y_train)  
model.intercept_
```


```
Out[13]: 4.440892098500626e-14
```

```
In [14]: model.coef_
```

```
Out[14]: array([ 7.33711071e-17, -3.50414142e-16, -2.27913527e-15,  6.36643  
516e-16,  
        -1.63064007e-16, -4.69906630e-14,  1.41813644e-16,  6.84944  
722e-17,  
        7.96888597e-18,  1.00000000e+00])
```

```
In [15]: prediction=model.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x7f8eca903880>
```



```
In [16]: model.score(x_test,y_test)
```

```
Out[16]: 1.0
```

```
In [ ]:
```

